

Ch. 6 – Point Estimation

Parameters and Estimators

A **parameter** is a constant describing a distribution:

- In a normal distribution, the mean μ and variance σ^2 are parameters.
- In an exponential distribution, the rate λ is a parameter.
- We will often use the symbol θ to represent a parameter generically.

An **estimator** is a random variable which is used to estimate a parameter.

- Given a random sample X_1, \dots, X_n , the sample mean $\bar{X} = \frac{1}{n}(X_1 + \dots + X_n)$ is an estimator for the mean μ .
- The sample variance $S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$ is an estimator for the variance σ^2 .
- We will often use the symbol $\hat{\theta}$ to represent an estimator generically.

Unbiased Estimators

An estimator $\hat{\theta}$ is an **unbiased** estimator for θ if $E(\hat{\theta}) = \theta$.

- If X is a binomial random variable, then the **sample proportion** $\hat{p} = \frac{X}{n}$ is an unbiased estimator of p :

$$E(\hat{p}) = E\left(\frac{X}{n}\right) = \frac{1}{n}E(X) = \frac{1}{n} \cdot np = p$$

- If X_1, \dots, X_n is a random sample from a distribution with mean μ , the sample mean \bar{X} is an unbiased estimator for μ :

$$\begin{aligned} E(\bar{X}) &= E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n} E\left(\sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} \sum_{i=1}^n E(X_i) = \frac{1}{n} \sum_{i=1}^n \mu = \frac{1}{n} \cdot n\mu = \mu \end{aligned}$$

Example: Uniform Distribution

Given a random sample X_1, \dots, X_n from a uniform distribution on $[0, \theta]$, how do we estimate θ ?

- The mean μ is the midpoint of the interval, $\mu = \theta/2$.
Therefore, $\theta = 2\mu$. Since we can estimate μ with \bar{X} , we can estimate θ with $\hat{\theta} = 2\bar{X}$.
- Each of the observations X_1, \dots, X_n will be less than θ , and if n is large we expect one of them to be close to θ . So we may estimate θ using the maximum: $\hat{\theta} = \max\{X_1, \dots, X_n\}$.

Example: Uniform Distribution

We gave two possible estimators for the parameter θ of a uniform distribution on $[0, \theta]$:

$$\hat{\theta}_1 = 2\bar{X}$$

$$\hat{\theta}_2 = \max\{X_1, \dots, X_n\}$$

Question: Are these estimators unbiased?

We may calculate the expected value of $\hat{\theta}_1$:

$$E(\hat{\theta}_1) = E(2\bar{X}) = 2E(\bar{X}) = 2\mu = \theta$$

Since $E(\hat{\theta}_1) = \theta$, this means that $\hat{\theta}_1$ is an unbiased estimator for θ .

Example: Uniform Distribution

Since $\hat{\theta}_2 = \max\{X_1, \dots, X_n\}$ is always less than θ , intuition suggests that $\hat{\theta}_2$ will underestimate θ and hence must be biased.

Now we will calculate $E(\hat{\theta}_2)$. For $0 \leq x \leq \theta$, the cdf $F(x)$ of $\hat{\theta}_2$ is

$$\begin{aligned} F(x) &= P(\hat{\theta}_2 \leq x) = P(\max\{X_1, \dots, X_n\} \leq x) \\ &= P(X_1 \leq x, X_2 \leq x, \dots, X_n \leq x) \\ &= P(X_1 \leq x)P(X_2 \leq x) \cdots P(X_n \leq x) = (x/\theta)^n \end{aligned}$$

Therefore the pdf of $\hat{\theta}_2$ is, for $0 \leq x \leq \theta$,

$$f(x) = F'(x) = \frac{d}{dx}(x/\theta)^n = nx^{n-1}/\theta^n$$

So the expected value of $\hat{\theta}_2$ is

$$E(\hat{\theta}_2) = \int_{-\infty}^{\infty} xf(x) dx = \int_0^{\theta} x \cdot \frac{nx^{n-1}}{\theta^n} dx = \frac{n}{n+1}\theta$$

Example: Uniform Distribution

We considered two estimators for the parameter θ of a uniform distribution on $[0, \theta]$,

$$\hat{\theta}_1 = 2\bar{X}$$

$$\hat{\theta}_2 = \max\{X_1, \dots, X_n\}$$

We found that $\hat{\theta}_1$ was unbiased but that $\hat{\theta}_2$ was biased. However, it is easy to modify $\hat{\theta}_2$ to produce an unbiased estimator $\hat{\theta}_3$:

$$\hat{\theta}_3 = \frac{n+1}{n}\hat{\theta}_2 = \frac{n+1}{n} \max\{X_1, \dots, X_n\}$$

Since $E(\hat{\theta}_2) = \frac{n}{n+1}\theta$, it follows that

$$E(\hat{\theta}_3) = E\left(\frac{n+1}{n}\hat{\theta}_2\right) = \frac{n+1}{n}E(\hat{\theta}_2) = \frac{n+1}{n} \cdot \frac{n}{n+1}\theta = \theta$$

so $\hat{\theta}_3$ is in fact an unbiased estimator for θ .

Example: Uniform Distribution

We now have two unbiased estimators for the parameter θ of a uniform distribution on $[0, \theta]$:

$$\hat{\theta}_1 = 2\bar{X}$$

$$\hat{\theta}_3 = \frac{n+1}{n} \max\{X_1, \dots, X_n\}$$

Question: Which of these estimators is better?

To answer this, we need a measure of how good an estimator is. One commonly used such measure is the *variance* of the estimator.

We can calculate the variance of $\hat{\theta}_1$:

$$V(\hat{\theta}_1) = V(2\bar{X}) = 4V(\bar{X}) = \frac{4}{n}V(X_1) = \frac{4}{n} \cdot \frac{\theta^2}{12} = \frac{\theta^2}{3n}$$

Example: Uniform Distribution

Recall the pdf of $\hat{\theta}_2$ is $f(x) = nx^{n-1}/\theta^n$. Therefore,

$$E(\hat{\theta}_2^2) = \int_0^\infty x^2 f(x) dx = \int_0^\infty nx^{n+1}/\theta^n dx = \frac{n}{n+2}\theta^2$$

$$\begin{aligned} V(\hat{\theta}_2) &= E(\hat{\theta}_2^2) - [E(\hat{\theta}_2)]^2 = \frac{n}{n+2}\theta^2 - \left(\frac{n}{n+1}\theta\right)^2 \\ &= \left(\frac{n(n+1)^2 - n^2(n+2)}{(n+1)^2(n+2)}\right)\theta^2 = \frac{n}{(n+1)^2(n+2)}\theta^2 \end{aligned}$$

$$\begin{aligned} V(\hat{\theta}_3) &= V\left(\frac{n+1}{n}\hat{\theta}_2\right) = \left(\frac{n+1}{n}\right)^2 V(\hat{\theta}_2) \\ &= \left(\frac{n+1}{n}\right)^2 \cdot \frac{n}{(n+1)^2(n+2)}\theta^2 = \frac{\theta^2}{n(n+2)} \end{aligned}$$

Example: Uniform Distribution

We considered two unbiased estimators for the parameter θ of a uniform distribution on $[0, \theta]$:

$$\hat{\theta}_1 = 2\bar{X}$$

$$\hat{\theta}_3 = \frac{n+1}{n} \max\{X_1, \dots, X_n\}$$

We calculated that their variances were $V(\hat{\theta}_1) = \frac{\theta^2}{3n}$ and $V(\hat{\theta}_3) = \frac{\theta^2}{n(n+2)}$. Therefore, for $n > 1$, $\hat{\theta}_3$ has a smaller variance.

More advanced statistical theory can be used to show that in fact $\hat{\theta}_3$ is a **minimum variance unbiased estimator**: it has a smaller variance than any other unbiased estimator.

Unbiasedness of Sample Variance

Given a distribution with variance σ^2 , the sample variance $S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$ is an unbiased estimator for σ^2 .

$$\begin{aligned} E(S^2) &= E \left[\frac{1}{n-1} \left(\sum_{i=1}^n X_i^2 - n\bar{X}^2 \right) \right] \\ &= \frac{1}{n-1} \left[\sum_{i=1}^n E(X_i^2) - nE(\bar{X}^2) \right] \\ &= \frac{1}{n-1} \left[\sum_{i=1}^n (V(X_i) + [E(X_i)]^2) - n(V(\bar{X}) + [E(\bar{X})]^2) \right] \\ &= \frac{1}{n-1} \left[\sum_{i=1}^n (\sigma^2 + \mu^2) - n\left(\frac{\sigma^2}{n} + \mu^2\right) \right] = \sigma^2 \end{aligned}$$

Maximum Likelihood Estimation

Suppose a Poisson process has unknown rate λ . We observe the process for 2 hours, and the number of events which occur is $X = 40$. What value of λ gives the largest probability $P(X = 40)$?

X is a Poisson random variable with mean 2λ , so

$$P(X = 40) = \frac{e^{-2\lambda}(2\lambda)^{40}}{40!}$$

The function $L(\lambda) = \frac{e^{-2\lambda}(2\lambda)^{40}}{40!}$ here is called the **likelihood function**. To maximize $L(\lambda)$, we take the derivative:

$$L'(\lambda) = \frac{1}{40!}(e^{-2\lambda} \cdot 40(2\lambda)^{39} \cdot 2 - 2e^{-2\lambda}(2\lambda)^{40})$$

Setting this equal to zero gives $40 - 2\lambda = 0$, so $\lambda = 20$. This is called the **maximum likelihood estimate** of λ .

Maximum Likelihood Estimation

In the previous problem, the likelihood function was given by

$$L(\lambda) = \frac{e^{-2\lambda}(2\lambda)^{40}}{40!}$$

The **log-likelihood function** is

$$\ell(\lambda) = \ln(L(\lambda)) = -2\lambda + 40 \ln(2\lambda) - \ln(40!)$$

The value of λ which maximizes $\ell(\lambda)$ will also maximize $L(\lambda)$. The log-likelihood function is often easier to differentiate than the likelihood function. In this case,

$$\ell'(\lambda) = -2 + \frac{40}{\lambda}$$

Setting this equal to zero gives $\lambda = 20$ as before.

The log-likelihood function was found to be

$$l(\lambda) = \ln(L(\lambda)) = -2\lambda + 40 \ln(2\lambda) - \ln(40!)$$

The maximum of log-likelihood function may be found by setting the derivative equal to zero:

$$l'(\lambda) = -2 + \frac{40}{2\lambda} \cdot 2 = -2 + \frac{40}{\lambda} = 0$$

Solving for λ gives $\lambda = 20$. This is the **maximum likelihood estimate** for λ .

Maximum Likelihood Estimation

In general, given a random variable X with pmf or pdf $f(x; \theta)$ depending on a parameter θ , the **maximum likelihood estimator** (mle) is the value $\hat{\theta}$ of θ that maximizes $f(X; \theta)$.

More generally, given a random sample X_1, \dots, X_n , the mle is the value $\hat{\theta}$ that maximizes the joint pmf or joint pdf of X_1, \dots, X_n :
$$f(X_1, \dots, X_n; \theta) = f(X_1; \theta)f(X_2; \theta) \cdots f(X_n; \theta).$$

Under fairly general conditions, the maximum likelihood estimator $\hat{\theta}$ satisfies some desirable statistical properties:

- $\hat{\theta}$ exists and is unique.
- $\hat{\theta}$ is *asymptotically unbiased*: for large sample sizes, it is practically an unbiased estimator.
- $\hat{\theta}$ is *asymptotically efficient*: for large sample sizes, it approximately achieves the smallest possible variance for an unbiased estimator.
- $\hat{\theta}$ is *asymptotically normal*: for large sample sizes, it has approximately a normal distribution.

Example

Given a random sample X_1, \dots, X_n from an exponential distribution with rate λ , find the mle for λ .

The pdf of each X_i is $f(x; \lambda) = \lambda e^{-\lambda x}$. Since X_1, \dots, X_n are independent, the likelihood function is

$$L(\lambda) = f(x_1, \dots, x_n; \lambda) = \prod_{i=1}^n \lambda e^{-\lambda x_i} = \lambda^n e^{-\lambda \sum_{i=1}^n x_i}$$

The log-likelihood is therefore

$$\ell(\lambda) = n \ln \lambda - \lambda \sum_{i=1}^n x_i$$

Setting the derivative equal to zero gives $\frac{n}{\lambda} - \sum_{i=1}^n x_i = 0$. Solving for λ , we find the mle:

$$\hat{\lambda} = \frac{n}{\sum_{i=1}^n X_i} = 1/\bar{X}$$

Example

Given an observation of an exponential random variable X with unknown rate λ , find the mle for λ .

The pdf of each X is $f(x) = \lambda e^{-\lambda x}$. The log-likelihood is therefore

$$\ell(\lambda) = \ln \lambda - \lambda x$$

Setting the derivative equal to zero gives $\frac{1}{\lambda} - x = 0$. Solving for λ , we find the mle:

$$\hat{\lambda} = \frac{1}{X}$$

Invariance Principle for MLE

Let $h(\theta)$ be a function of θ . If $\hat{\theta}$ is the mle for θ , then $h(\hat{\theta})$ is the mle for $h(\theta)$.

Example: Given an observation of an exponential random variable X with unknown rate λ , we found that the mle for λ was

$$\hat{\lambda} = \frac{1}{X}$$

Therefore, the mle for $\mu = 1/\lambda$ is

$$\hat{\mu} = 1/\hat{\lambda} = \frac{1}{1/X} = X$$

Thus, the mle for μ is simply the observed value X . It can be shown that this in fact is the minimum variance unbiased estimator for μ .

MLE of Several Parameters

There are many families of distributions where we would want to estimate two parameters at once. For example,

- For a normal distribution, we want to estimate μ and σ .
- For a gamma distribution, we want to estimate k and λ .
- For a Weibull distribution, we want to estimate α and β .

The definition of mle extends in an natural way to distributions with several parameters:

Given a random sample X_1, \dots, X_n from a distribution with several parameters $\theta_1, \dots, \theta_m$, the mle of $(\theta_1, \dots, \theta_m)$ is the combination of values $(\hat{\theta}_1, \dots, \hat{\theta}_m)$ maximizing the joint pmf or joint pdf $f(x_1, \dots, x_n; \theta_1, \dots, \theta_m)$ of X_1, \dots, X_n .

Estimating Parameters of a Normal Distribution

Given a random sample X_1, \dots, X_n from a normal distribution:

The MLE for μ is the sample mean \bar{X} , and this is also the minimum variance unbiased estimator.

This means that alternative estimators for μ , such as the sample median \tilde{X} , must have a greater variance. (In fact, $V(\bar{X}) = \frac{\sigma^2}{n}$, while $V(\tilde{X}) \approx \frac{\pi}{2} \cdot \frac{\sigma^2}{n}$ for large n .)

The MLE for σ^2 is $\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$. However, this estimator is biased; the minimum variance unbiased estimator is the sample variance $S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$.